

Oracle Database – Memory and Caching Technologies

#oracledatabase #oraclecloud #autonomousdatabase #memory #caching #tables #cursors
#resultcache #sql #plsql #19c #23ai #truecache #inmemory #bufferpools #tipps
#performance #advisor #fastingest #fastlookup #memoptimize

Ulrike Schwinn

Oracle Data Management Expert



Videos, Blogs, GitHub from Oracle EMEA Experts and Specialists



YouTube: 23ai Playlist OraTech



GitHub: converged database



Blogs: blogs.oracle.com/coretec

Safe harbor statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, timing, and pricing of any features or functionality described for Oracle's products may change and remains at the sole discretion of Oracle Corporation.

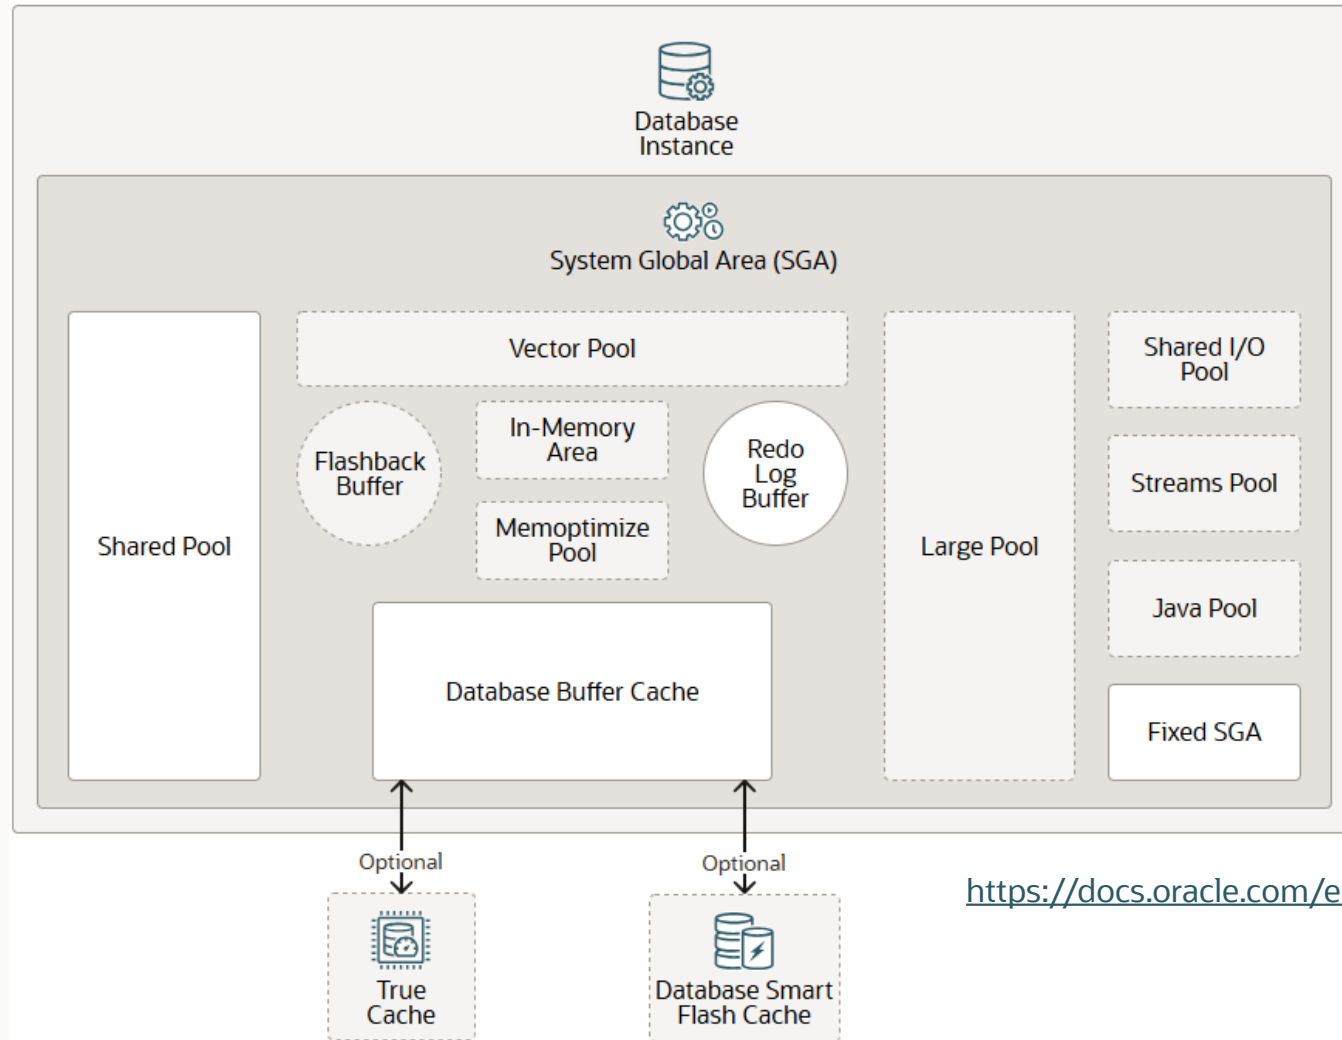


Some Performance Aspects

- When to check performance?
 - Loss in performance/Degradation/Complaints
 - New Versions/Feature Testing
 - Migration/Upgrade
- Focus: on statements/application code with „bad“ performance
- Know your **available access structure and design**
 - Indexes, IOT, MAVs, Views, Partitioning, data types, ...
- Make sure **statistics** are up-to-date
- Dont's
 - Row by Row Processing
- Do's
 - Bulk Loads, Direct Loads, Parallelization
 - Result Cache for compute intense queries
 - Object Caching/Warming up
 - Memory optimizing techniques such as compression, partitioning etc.
 - Use Advisory and monitoring framework and automations (if possible)



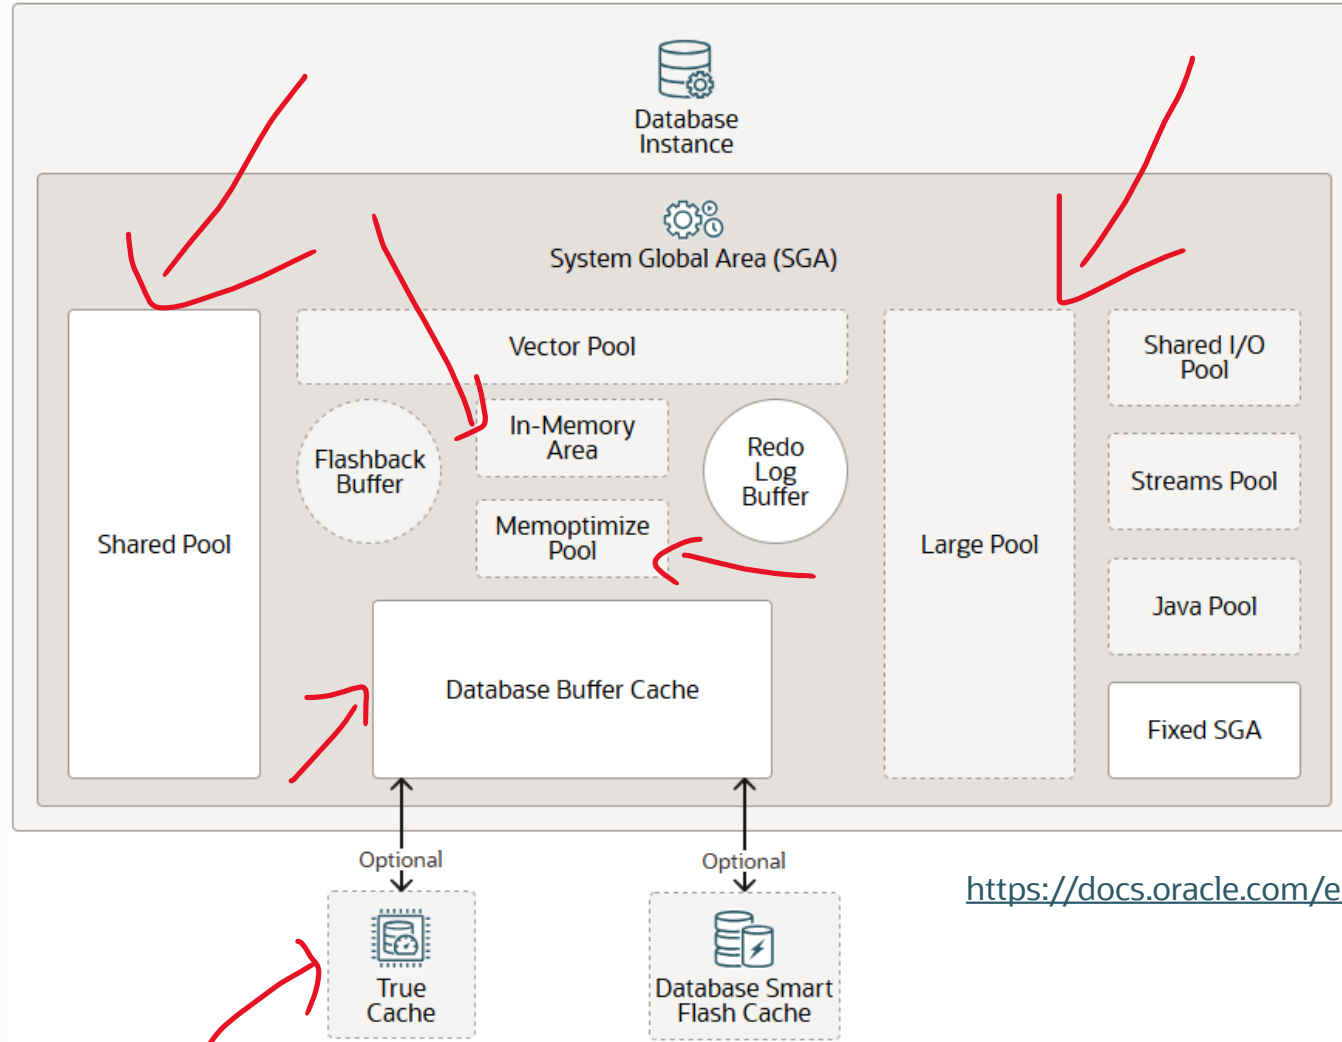
Oracle Database 23ai Technical Architecture – System Global Area



https://docs.oracle.com/en/database/oracle/oracle-database/23/dbiad/db_sga.html




Oracle Database 23ai Technical Architecture – System Global Area




https://docs.oracle.com/en/database/oracle/oracle-database/23/dbiad/db_sga.html




Database Buffer



Result Cache



Database Smart Flash Cache Buffer




Flash Cache Level 2

Oracle Database Memory and Caching

23^{ai}


True Cache



High Speed Data Ingestion MemOptimized RowStore



Database In-Memory



Exadata with HCC
Smart Scan
Storage Index
Smart Flash Cache

Oracle TimesTen
In-Memory Database
memory-optimized relational database
deployed in the middle-tier

Coherence
Data Grid
in-memory distributed data grid
solution for clustered applications

Database Buffer

SGA, LRU
Keep&Recycle Pools
Table data caching
PL/SQL& cursor caching

Result Cache


Result Cache Area
in Shared Pool
for compute intense
queries and PL/SQL
function results

High Speed
Data Ingestion
MemOptimized RowStore




Oracle Database Memory and Caching

Database Smart Flash Cache Buffer




Flash Cache Level 2

Database In-Memory



23^{ai}

True Cache



List SGA Components

- `v$sga_dynamic_components` summarizes information based on all completed SGA resize operations.

COMPONENT	CURRENT_SIZE
-----	-----
shared pool	2.9595E+10
large pool	4362076160
java pool	1610612736
streams pool	268435456
unified pga pool	0
memoptimize buffer cache	536870912
DEFAULT buffer cache	9.9724E+10
KEEP buffer cache	0
RECYCLE buffer cache	0
DEFAULT 2K buffer cache	0
...	
In-Memory Area	0

See also [Documentation](#)

- Use Memory Advisor with `DB_CACHE_ADVICE` set to ON
 - Enables statistics gathering used for predicting behavior with different cache sizes through the `v$db_cache_advice` performance view.



Usage of V\$DB_CACHE_ADVICE

- Set **DB_CACHE_ADVICE** ON. If STATISTICS_LEVEL is TYPICAL or ALL, then default is already ON.
- Query V\$DB_CACHE_ADVICE - here for default buffer pools.

```
SELECT size_for_estimate size, buffers_for_estimate buffers, estd_physical_read_factor estd_read,
estd_physical_reads estd_write
FROM V$DB_CACHE_ADVICE
WHERE name = 'DEFAULT' AND block_size =
(SELECT value FROM V$PARAMETER WHERE name = 'db_block_size');
```

SIZE	BUFFERS	ESTD_READ_FACTOR	ESTD_WRITE	
-----	-----	-----	-----	
30	3,802	18.70	192,317,943	<== 10% of Current Size
60	7,604	12.83	131,949,536	
...				
212	26,614	1.74	17,850,847	<== Current Size
243	30,416	1.33	13,720,149	
273	34,218	1.13	11,583,180	
304	38,020	1.00	10,282,475	
334	41,822	.93	9,515,878	
364	45,624	.87	8,909,026	
395	49,426	.83	8,495,039	


See also [Documentation](#)



Buffer Cache Pools: KEEP and RECYCLE Pools

- Manually configure separate buffer pools that either **keep data** in the buffer cache or **make the buffers available for new data** immediately after using the data blocks.
- Pool configuration with `DB_KEEP_CACHE_SIZE`, `DB_RECYCLE_CACHE_SIZE`
- Assign schema object with ALTER/CREATE command

```
ALTER TABLE ... STORAGE (buffer_pool keep)
ALTER INDEX ... STORAGE (buffer_pool keep)
ALTER TABLE ... MODIFY LOB (lobcol) (STORAGE (buffer_pool keep))
```
- When to consider:
 - Objects in KEEP pool are HOT and should not be aged out.
 - Objects in RECYCLE pool should not consume unnecessary space in the cache.



See also [Documentation](#)

Are database objects cached?

- Check execution plans, statistics, or V\$ views such as V\$BH
- **v\$dbh** displays the status and number of pings for every buffer in the SGA.
- Example:

```
SELECT o.object_name, o.object_type, o.owner, COUNT(*) NUMBER_OF_BLOCKS
FROM dba_objects o, v$dbh bh
WHERE o.data_object_id = bh.objd
AND o.owner in ('SH')
GROUP BY o.object_name, o.owner, o.object_type
ORDER BY COUNT(*) desc;
```

OBJECT_NAME	OBJECT_TYPE	OWNER	NUMBER_OF_BLOCKS
-----	-----	-----	-----
CUSTOMERS	TABLE	SH	1673
SALES	TABLE PARTIT	SH	233
SALES_PROMO_BIX	INDEX PARTITION	SH	58
...			

See also [Documentation](#)



Objects Caching/Pre-Loading

- Initiate object scans over tables, indexes , lob segments etc.

```
SELECT /*+ FULL(T1) */ sum(numeric_column), min(txt_column) FROM tabelle T1;  
SELECT /*+ FULL(T2) */ dbms_lob.getlength(lob_column) FROM tabelle T2;
```

- Keep in mind: Short table scans are scans performed on tables under a certain size threshold. The definition of a small table is the maximum of 2% of the buffer cache or 20, whichever is bigger.
- Big tables (= segments) are not cached when read with a full scan.
 - By default, a big table (=segment) is typically a table (=segment) bigger than 2% of the buffer cache.
 - This behavior can be controlled by setting the "_small_table_threshold" parameter, that specifies the threshold in database blocks for deciding if a segment should be cached.
 - Modifiable with ALTER SESSION to get a bigger size

See also [Documentation](#)



Check hidden parameters

```
SELECT a.ksppinm "Parameter", b.KSPSTDF "Default", b.kspstvl "Session", c.kspstvl "Instance",
       decode(bitand(a.ksppiflg/256,1),1,'TRUE','FALSE') IS_SESSION_MOD,
       decode(bitand(a.ksppiflg/65536,3),1,'IMMEDIATE',2,'DEFERRED',3,'IMMEDIATE','FALSE') IS_SYSTEM_MOD
FROM x$ksppi a,x$ksppcv b, x$ksppsv c
WHERE a.indx = b.indx AND a.indx = c.indx AND UPPER(ksppinm) like UPPER('%&name%');
```

```
Enter value for name: _small_table
old 5: AND UPPER(ksppinm) like UPPER('%&name%')
new 5: AND UPPER(ksppinm) like UPPER('%_small_table%')
```

Parameter	Default	Session	Instance	IS_SESSION_MOD	IS_SYSTEM_MOD
-----	-----	-----	-----	-----	-----
_small_table_threshold	TRUE	220314	220314	TRUE	DEFERRED
...					

[How To Query And Change The Oracle Hidden Parameters In Oracle 10g and Later \(Doc ID 315631.1\)](#)



PL/SQL Object Caching

- To improve performance and avoid reparsing, prevent SQL or PL/SQL areas from aging out of the shared pool.
- `dbms_shared_pool` enables you to mark them for keeping or not-keeping , to purge, to markhot or unmarkhot

- For frequently used triggers, procedures and tables, also for sequences, use KEEP

```
execute dbms_shared_pool.keep(name=> 'SH.TESTP1', flag => 'P');
execute dbms_shared_pool.keep(name=> '4158E358 2329752635', flag=> 'C');
```

- For „Library cache: mutex X” as a top events, consider MARKHOT

```
execute dbms_shared_pool.markhot('SH.TESTP1', 'P');
```

Top 5 Timed Foreground Events

Event	Waits	Time(s)	Avg wait (ms)	% DB time	Wait Class
library cache: mutex X	53,448,946	312,984	6	74.18	Concurrency

- Monitoring in v\$db_object_cache

```
SQL> select name, kept, status from v$db_object_cache where name like 'TESTP1';
```

NAME	KEPT	STATUS
-----	-----	-----
TESTP1	YES	VALID


See also [Documentation](#)



SQL Query and PL/SQL Function Result Cache

Server Result Cache contains:

- **SQL query result cache** stores the results of queries and query fragments.
- **PL/SQL function result cache** stores function result sets.



Server Result Cache	On database startup, Oracle Database allocates memory to the server result cache in the shared pool.
SQL Query Result Cache	<ul style="list-style-type: none">• Result cache hints at the application level <code>SELECT /*+ RESULT_CACHE */ ...</code>• Table annotations affect the entire query <code>CREATE TABLE sales (...) RESULT_CACHE (MODE DEFAULT)</code>
PL/SQL Function Result Cache	<ul style="list-style-type: none">• Needs RESULT_CACHE clause

When to consider:

- Predictable queries/repeating queries
- Long-running queries with expensive calculations
- Functions that are invoked frequently but depend on information that changes infrequently

See also [Documentation](#)



Parameters

RESULT_CACHE_MAX_RESULT		5%
RESULT_CACHE_MAX_SIZE		% of RESULT_CACHE_MAX_SIZE any single result can use derived from memory components (0 means disabled)
RESULT_CACHE_MODE		<u>MANUAL</u> MANUAL_TEMP FORCE FORCE_TEMP
RESULT_CACHE_REMOTE_EXPIRATION		minutes that a result using a remote object is allowed to remain valid.
RESULT_CACHE_MAX_TEMP_RESULT	21^c	5 % % of RESULT_CACHE_MAX_TEMP_SIZE
RESULT_CACHE_MAX_TEMP_SIZE	21^c	RESULT_CACHE_MAX_SIZE * 10 maximum amount of temporary tablespace that can be consumed by the result cache.
RESULT_CACHE_INTEGRITY	23^{ai}	ENFORCED <u>TRUSTED</u> (*)
RESULT_CACHE_EXECUTION_THRESHOLD	21^c	2 maximum number of times a PL/SQL function can be executed before its result is stored in the result cache.
RESULT_CACHE_AUTO_BLOCKLIST	23^{ai}	<u>ON</u> OFF to enable or disable adaptive result cache object exclusion

See also [Documentation](#)



Usage in Autonomous Database and Exadata

NAME	VALUE
result_cache_integrity	ENFORCED
result_cache_mode	FORCE
result_cache_max_size	5242880
result_cache_max_result	1
result_cache_remote_expiration	0
result_cache_execution_threshold	2
result_cache_max_temp_size	536870912
result_cache_max_temp_result	5
result_cache_auto_blocklist	ON

```
SQL> select name, value from v$parameter  
       where name like 'result%';
```

NAME	VALUE
-----	-----
result_cache_integrity	TRUSTED
result_cache_mode	MANUAL
result_cache_max_size	4194304
result_cache_max_result	5
result_cache_remote_expiration	0
result_cache_execution_threshold	2
result_cache_max_temp_size	41943040
result_cache_max_temp_result	5
result_cache_auto_blocklist	ON

Result Cache: Query Hints

```
SELECT /*+ result_cache */ a.department_id "Department",
      a.num_emp/b.total_count "%_Employees",
      a.sal_sum/b.total_sal "%_Salary"
FROM
  (SELECT /*+ result_cache */ department_id, COUNT(*) num_emp, SUM(salary) sal_sum
   FROM employees
   GROUP BY department_id) a,
  (SELECT /*+ result_cache */ COUNT(*) total_count, SUM(salary) total_sal
   FROM employees) b
ORDER BY a.department_id;
```


Results Script Output Explain Autotrace DBMS Output OWA Output

Operation	Optimizer	Cost	Cardinality	Bytes	Partition Start	Partition Stop	Partition Id
SELECT STATEMENT	ALL_ROWS	7	11	715			
RESULT CACHE 7v838kf3k3zkg0yqpabtt131dt							
SORT(ORDER BY)		7	11	715			
NESTED LOOPS		7	11	715			
VIEW		3	1	26			
RESULT CACHE a6jgaqk8d0wzq48umutf26n89h							
SORT(AGGREGATE)			1	4			
TABLE ACCESS(FULL) HR.EMPLOYEES ANALYZED		3	107	428			
VIEW		4	11	429			
RESULT CACHE a2ujsg3da29mhc8rhhmga0ygz7							
SORT(GROUP BY)		4	11	77			
TABLE ACCESS(FULL) HR.EMPLOYEES ANALYZED		3	107	749			

See also [Documentation](#)



Enabling PL/SQL Result Cache for Functions



- To make a function result-cached, you must include the `RESULT_CACHE` clause in the function declaration and definition.

```
CREATE OR REPLACE FUNCTION get_date (p_id number, p_format varchar2) RETURN varchar2
RESULT_CACHE
IS
    v_date date;
BEGIN
    select hiredate into v_date
    from emp where empno = p_id;

    RETURN to_char(v_date, p_format);
END;
```

- 21^c** • `RESULT_CACHE_EXECUTION_THRESHOLD`: max number of times a PL/SQL function can be executed before its result is stored in the result cache.

See also [Documentation](#)

Administration

Use `dbms_result_cache.status | memory_report | flush | bypass`
and `v$views` such as

```
SQL> select name, type, row_count, status, invalidations, scan_count
      from v$result_cache_objects where scan_count!=0
```


NAME	TYPE	ROW_COUNT	STATUS	INVALIDATIONS	SCAN_COUNT

"SCOTT"."GET_DATE"::8."GET_DATE"#3a2b06cfa1322e42 #1	Result	1	Published	0	2
select /*+ result_cache */ c Result		1	Published	0	1
ount(*) from emp					


See also [Documentation](#)



Database Buffer



Result Cache




High Speed Data Ingestion
MemOptimized RowStore
Large Pool to buffer INSERTs
Memoptimize Pool for fast lookups




Oracle Database Memory and Caching

Database Smart Flash Cache Buffer




Flash Cache Level 2

Database In-Memory



23^{ai}

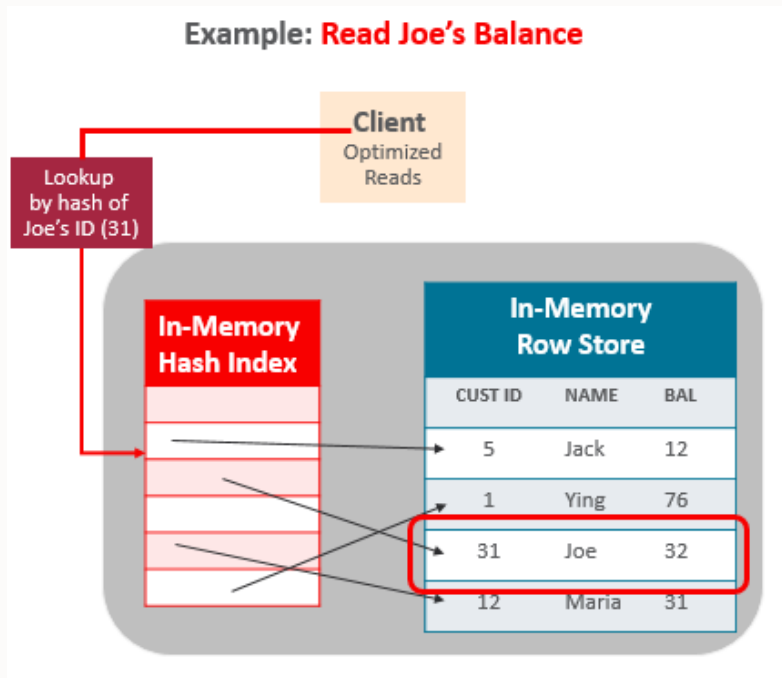
True Cache



Data Streaming: Fast Lookup and Fast Ingest

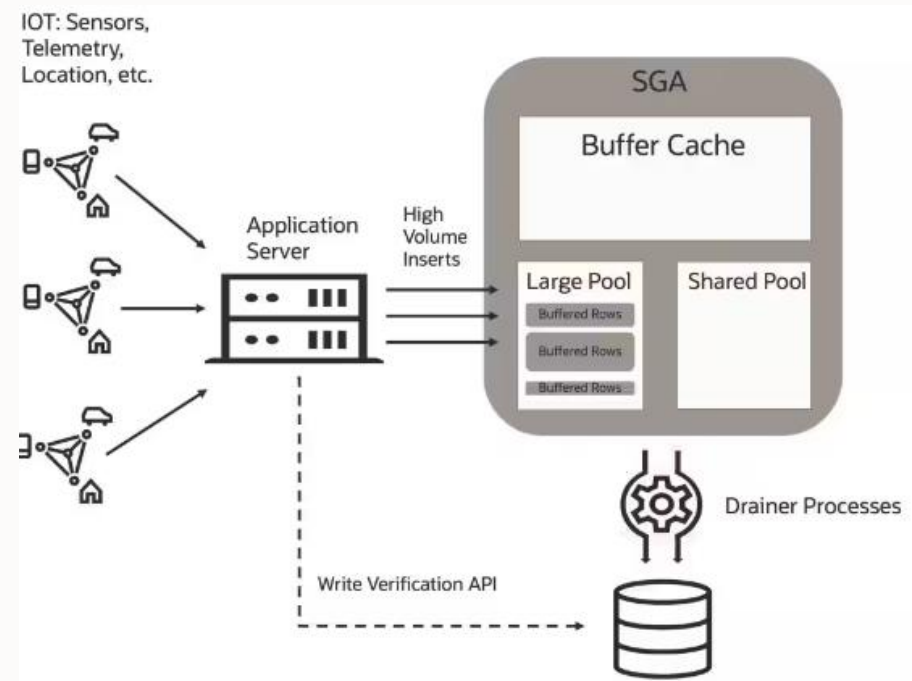
Fast Lookup

Scenarios: Data streaming applications like Internet of Things (IoT) **request/read** data for clients at a very high frequency.



Fast Ingest

Scenarios: IoT data from sensor, smart meter or even traffic camera might be collected and **written** to the database in high volumes for **later analysis**.



Included in Enterprise Edition from 19.12 onwards and in Database 23ai Free edition

Fast Lookup for high frequency queries - part 1

- Special **memoptimize pool** in SGA with buffer area and hash index
- Automatically managed
- For queries **with a single equality predicate on the primary key**

Workflow

1. Change MEMOPTIMIZE_POOL_SIZE in ROOT and restart
2. Add table attribute MEMOPTIMIZE FOR READ
3. Run DBMS_MEMOPTIMIZE.POPULATE
4. Query data!

```
-- in container ROOT set parameter and restart
SQL> alter system set MEMOPTIMIZE_POOL_SIZE = 500M SCOPE=SPFILE;
System altered.
-- in pdb
SQL> create table sh.sales_tab (sales_id    NUMBER(6) primary key,
                                prod_id     NUMBER(6) not null,
                                cust_id     NUMBER not null,
                                time_id     DATE not null,
                                quantity_sold NUMBER(3) not null,
                                amount_sold  NUMBER(10,2) not null);
```

Table created.

```
SQL> insert into sh.sales_tab
(sales_id,prod_id, cust_id, time_id, quantity_sold, amount_sold)
select rownum,prod_id,cust_id,time_id,quantity_sold,amount_sold
from sh.sales;
918843 rows created.
```

```
-- enable it for the data stored in SH.SALES_TAB.
SQL> alter table sh.sales_tab MEMOPTIMIZE FOR READ;
Table altered.
SQL> execute dbms_stats.gather_table_stats('SH','SALES_TAB');
PL/SQL procedure successfully completed.
```



Fast Lookup for high frequency queries – part 2

- For queries **with a single equality predicate on the primary key**

```
--Populate the rowstore with DBMS_MEMOPTIMIZE.POPULATE
SQL> execute dbms_memoptimize.populate(schema_name=>'SH', table_name=>'SALES_TAB');
PL/SQL procedure successfully completed.
```

```
SQL> select * from sh.sales_tab where sales_id=5;
```

```
...
SQL> select * from dbms_xplan.display_cursor();
PLAN_TABLE_OUTPUT
```

```
-----
SQL_ID  dfn2vxgz8vuyf, child number 1
-----
select * from sh.sales_tab where sales_id=5
Plan hash value: 1081214910
-----
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT				3 (100)	
1	TABLE ACCESS BY INDEX ROWID READ OPTIM	SALES_TAB	1	27	3 (0)	00:00:01
* 2	INDEX UNIQUE SCAN READ OPTIM	SYS_C0011070	1		2 (0)	00:00:01

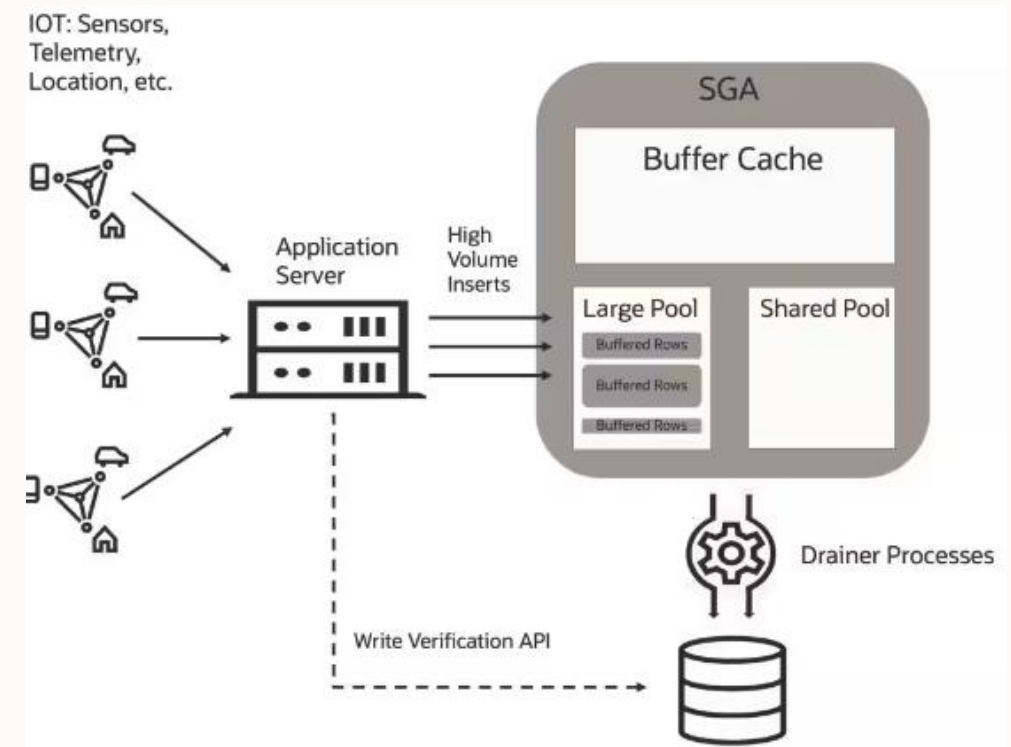
```
-----
```

```
Predicate Information (identified by operation id):
-----
  2 - access("SALES_ID"=5)
```



Fast Ingest for high-frequency data inserts

- Introduced in 19c
- Fast Ingest operations are very different from normal Oracle Database transaction processing.
- The normal transaction mechanisms are bypassed and the application needs to check the data was written.
 - COMMITs have no meaning.
 - There is no ability to rollback the inserts.
 - You cannot query data until it has been flushed to disk.
 - Index operations and constraint checking is done only when the data is written from the fast ingest area in the large pool to disk.
- Special packages like DBMS_MEMOPTIMIZE and DBMS_MEMOPTIMIZE_ADMIN support the technology.



Fast Ingest for high-frequency data inserts

- Uses the **large pool** to buffer the INSERTs before they are written to the hard disk

In **23^{ai}**

- MEMOPTIMIZE_WRITE_AREA_SIZE enables explicit allocation for the memoptimize write area in large pool.
- MEMOPTIMIZE_WRITES parameter avoids having to use the MEMOPTIMIZE_WRITE hint.

```
-- in container ROOT: the size of the ingest buffer in large pool
SQL> alter system set
      memoptimize_write_area_size=300M scope=spfile;
System altered.
```

```
-- restart the database
-- connect to the pdb
-- create a table for fast ingest
SQL> create table sh.sales_write_tab
      (sales_id      NUMBER(6) primary key,
       prod_id       NUMBER(6) not null,
       cust_id       NUMBER not null,
       time_id       DATE not null,
       quantity_sold NUMBER(3) not null,
       amount_sold   NUMBER(10,2) not null)
      segment creation immediate
      memoptimize for write;
```

Table created.



Demo: Fast Ingest

- DBMS_MEMOPTIMIZE.WRITE_END flushes fast ingest data from large pool.
- Fast Ingest supports now also
23^{ai} LOBs, Compression, Encryption, and Partitioning (interval, hash and sub-partition)

```
-- instruct the database to always use fast ingest for INSERTs
SQL> alter session set memoptimize_writes=on;
Session altered.
```

```
SQL> insert into sh.sales_write_tab
(sales_id,prod_id,cust_id,time_id,quantity_sold,amount_sold)
select rownum,prod_id,cust_id,time_id,quantity_sold,amount_sold
from sh.sales;
918843 rows created.
```

```
SQL> select count(*) from sh.sales_write_tab;
COUNT(*)
-----
113617
```

```
-- flush all the fast ingest data from large pool to disk
```

```
SQL> exec dbms_memoptimize.write_end;
PL/SQL procedure successfully completed.
```

```
SQL> select count(*) from sh.sales_write_tab;
COUNT(*)
-----
918843
```



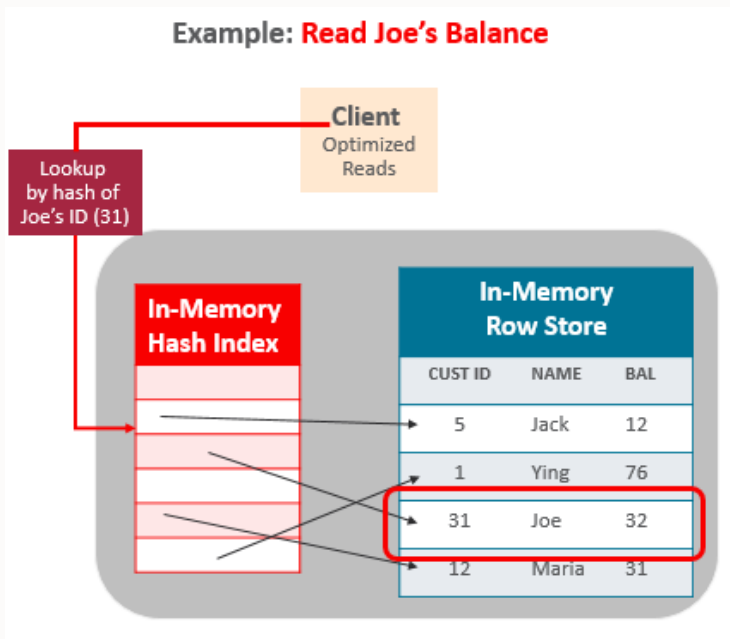
Fast Lookup and Fast Ingest in Oracle Database

Fast Lookup

Introduced in 18c

Uses special **memoptimize pool** in SGA with buffer area and hash index

For queries **with a single equality predicate on the primary key**

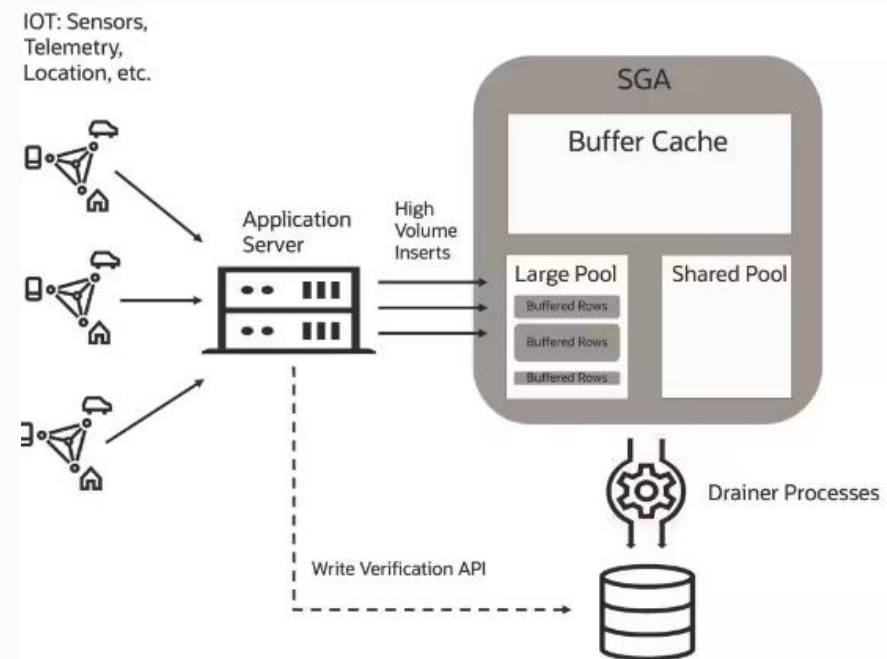


Fast Ingest

Introduced in 19c

Uses **large pool** to buffer the INSERTs before they are written to disk

In **23^{ai}**, improved memory handling and new supported objects and technology



[High Performance Data Streaming with Fast Lookup and Fast Ingest \(YouTube\)](#)

Further Readings

Documentation

- [Database Performance Tuning Guide 19c](#)
- [Database Performance Tuning Guide 23ai](#)
- [Database Concepts 23ai](#)
- [PL/SQL Packages and Types References 23ai](#)
- [Database Licensing Information User Manual](#)

Blogs, Videos & White Paper


- [Oracle Database 23ai Fast Ingest Enhancements](#)
- [New in Oracle Database 19c: Memoptimized Rowstore – Fast Ingest](#)
- [Best Practices For High Volume IoT workloads with Oracle Database 19c \(White Paper\)](#)
- [Fast Lookup with Memory Optimized Rowstore](#)
- [High Performance Data Streaming with Fast Lookup and Fast Ingest \(YouTube\)](#)

- **All Links**

<https://github.com/oracle-devrel/technology-engineering/tree/main/data-platform/core-converged-db/fast-ingest-lookup>



Database Buffer



Result Cache

Server Result Cache

SQL Query Result Cache


PL/SQL Function
Result Cache

High Speed
Data Ingestion
MemOptimized RowStore



Oracle Database Memory and Caching

Database Smart Flash Cache Buffer



Flash Cache Level 2


Database In-Memory

Row- and columnar
formats
for real-time analytics
and mixed workloads




23^{ai}

True Cache



Database In-Memory: Enables Dual Formatted Data




- **BOTH** row and column formats for same table
- Simultaneously active and transactionally consistent
- Analytics & reporting use new in-memory Column format
- OLTP uses proven row format
- No application changes required

[Oracle Database In-Memory Resource Page](#)


Improves All Aspects of Analytic Queries

Scans




- Speed of memory
- Scan and Filter only the needed Columns
- Vector Instructions

Joins



- Convert Star Joins into 10X Faster Column Scans
- Search large table for values that match small table

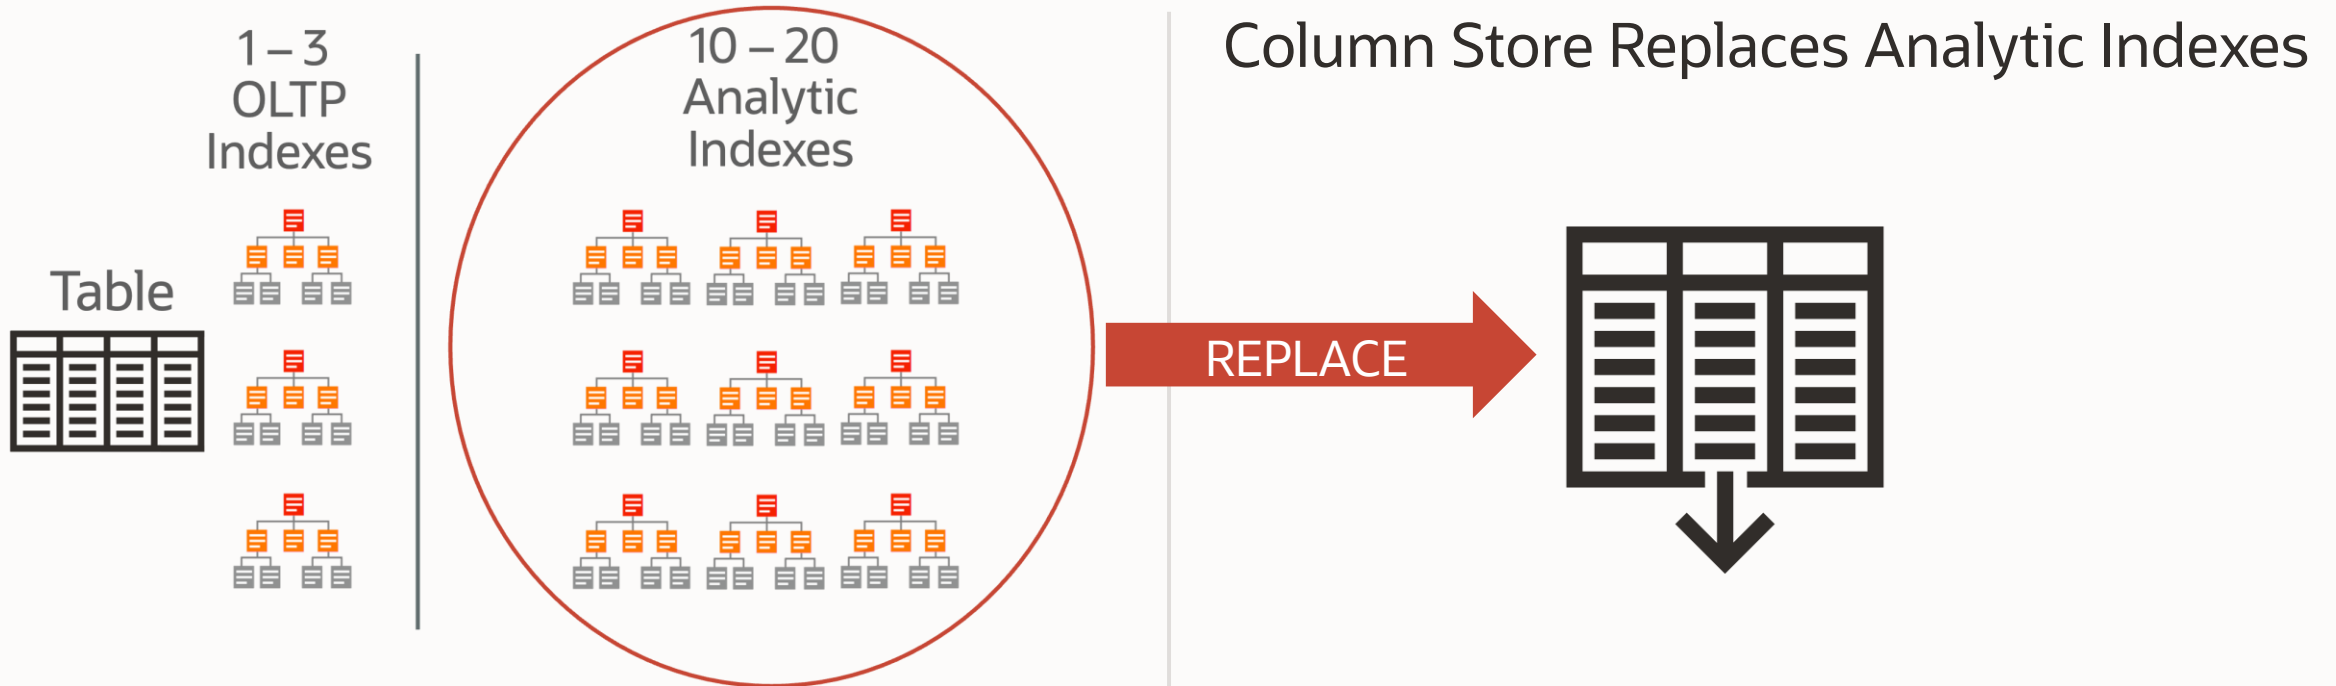
Reporting



- Create In-Memory Report Outline that is Populated during Fast Scan
- Runs Reports Instantly



Accelerates Mixed Workloads



- Inserting one row into a table requires updating 10-20 analytic indexes: **Slow!**
- Fast analytics only on indexed columns
- Analytic indexes **increase** database size

- Fast analytics on any column
- Column Store not persistent so updates are: Fast
- No analytic indexes: **Reduces** database size

Database In-Memory: Simple to Implement and Simple to Test

1. Configure Memory Capacity
`inmemory_size = XXX GB`
2. Configure tables or partitions to be in memory
`alter table | partition ...
inmemory;`
3. Later drop analytic indexes to speed up OLTP

Note: In 19.8, with `BASE_LEVEL` value for the `INMEMORY_FORCE`, `INMEMORY_SIZE` parameter can be set up to a value of 16GB without having to license the Database In-Memory option.

```
SQL> -- In-Memory Column Store query
SQL>
SQL> select  max(lo_ordtotalprice) most_expensive_order From LINEORDER;

MOST_EXPENSIVE_ORDER
-----
                    57346348
Elapsed: 00:00:00.01

| Id | Operation                                | Name      | Rows  | Bytes | Cost (%CPU)| Time      |
|----|-----|-----|-----|-----|-----|-----|-----|
|  0 | SELECT STATEMENT                        |           |      1 |      6 | 5401 (100)|           |
|  1 | SORT AGGREGATE                          |           |      1 |      6 |           |           |
|  2 | TABLE ACCESS INMEMORY FULL             | LINEORDER |    59M| 343M | 5401 (16) | 00:00:01 |

SQL> -- Buffer Cache query with the column store disabled via NO_INMEMORY hint
SQL>
SQL> select /*+ NO_INMEMORY */ max(lo_ordtotalprice) most_expensive_order From LINEORDER;

MOST_EXPENSIVE_ORDER
-----
                    57346348
Elapsed: 00:00:00.38

| Id | Operation                                | Name      | Rows  | Bytes | Cost (%CPU)| Time      |
|----|-----|-----|-----|-----|-----|-----|-----|
|  0 | SELECT STATEMENT                        |           |      1 |      6 | 123K(100)|           |
|  1 | SORT AGGREGATE                          |           |      1 |      6 |           |           |
|  2 | TABLE ACCESS FULL                      | LINEORDER |    59M| 343M | 123K (1)  | 00:00:05 |
```

• See also [“New In-Memory Eligibility Test”](#)



Further Readings

Documentation

- [Database In-Memory](#)
- [Database In-Memory Guide](#)

Blogs and technical briefs

- [DBIM Resources](#)
- [Oracle Database In-Memory](#)
- [When to Use Oracle Database In-Memory](#)
- [Oracle Database In-Memory Implementation Guidelines](#)

LiveLabs


- [Boost Analytics Performance with Oracle Database In-Memory](#)
- [Database In-Memory Advanced Features](#)

• All Links

<https://github.com/oracle-devrel/technology-engineering/tree/main/data-platform/core-converged-db/db-in-memory>



Database Buffer



Result Cache

Server Result Cache

SQL Query Result Cache


PL/SQL Function
Result Cache

High Speed
Data Ingestion
MemOptimized RowStore




Oracle Database Memory and Caching

Database Smart Flash Cache Buffer



Flash Cache Level 2

Database In-Memory



23^{ai}

True Cache
In-memory,
consistent, and
automatically
managed,
read-only SQL cache



What is Oracle True Cache ?

Oracle True Cache is a light-weight (nearly) disk-less Oracle database instance that is deployed as a cache


ANY SQL Query can be easily directed to the cache instead of the back-end database

- **Unique** transparent full-function data cache

Data requested by queries not found in True Cache is **automatically** retrieved from the back-end database

- Changes to data are **automatically** propagated to True Cache in **real-time**
- True Cache is **consistent** as of a point in time


[Further readings and links](#)



How To Use True Cache

The application can query data from True Cache in the following two modes:

1. The application maintains two connections:
 - **A read-only connection to True Cache** and **a read-write connection to the database** and uses the read-only connection for offloading queries to True Cache
2. The application maintains one connection.
 - The JDBC driver maintains two connections internally and does the read-write split under application control
 - The application uses `setReadOnly()` calls to mark some sections of code as “read-only”
 - This is an existing JDBC API which MySQL already uses for similar purposes
 - JDBC driver will send read-only queries to True Cache and writes and DDLs to the primary



True Cache Use Cases

True Cache can be deployed as a:



Mid-Tier Cache



Edge Cache



Cross-Region Cache



Cross-Cloud Cache

Further Readings

Documentation

- [True Cache on oracle.com](#)
- [True Cache FAQ](#)
- [Oracle True Cache User's Guide](#)
- [Java Support for True Cache](#)
- [Oracle True Cache Technical Architecture](#)

Blogs

- [Introducing Oracle True Cache: In-memory, consistent, and automatically managed SQL cache \(Oracle Database 23ai\)](#)
- [Getting started with True Cache in Oracle Database 23ai FREE \(Blog\)](#)
- [Oracle True Cache \(YouTube\)](#)

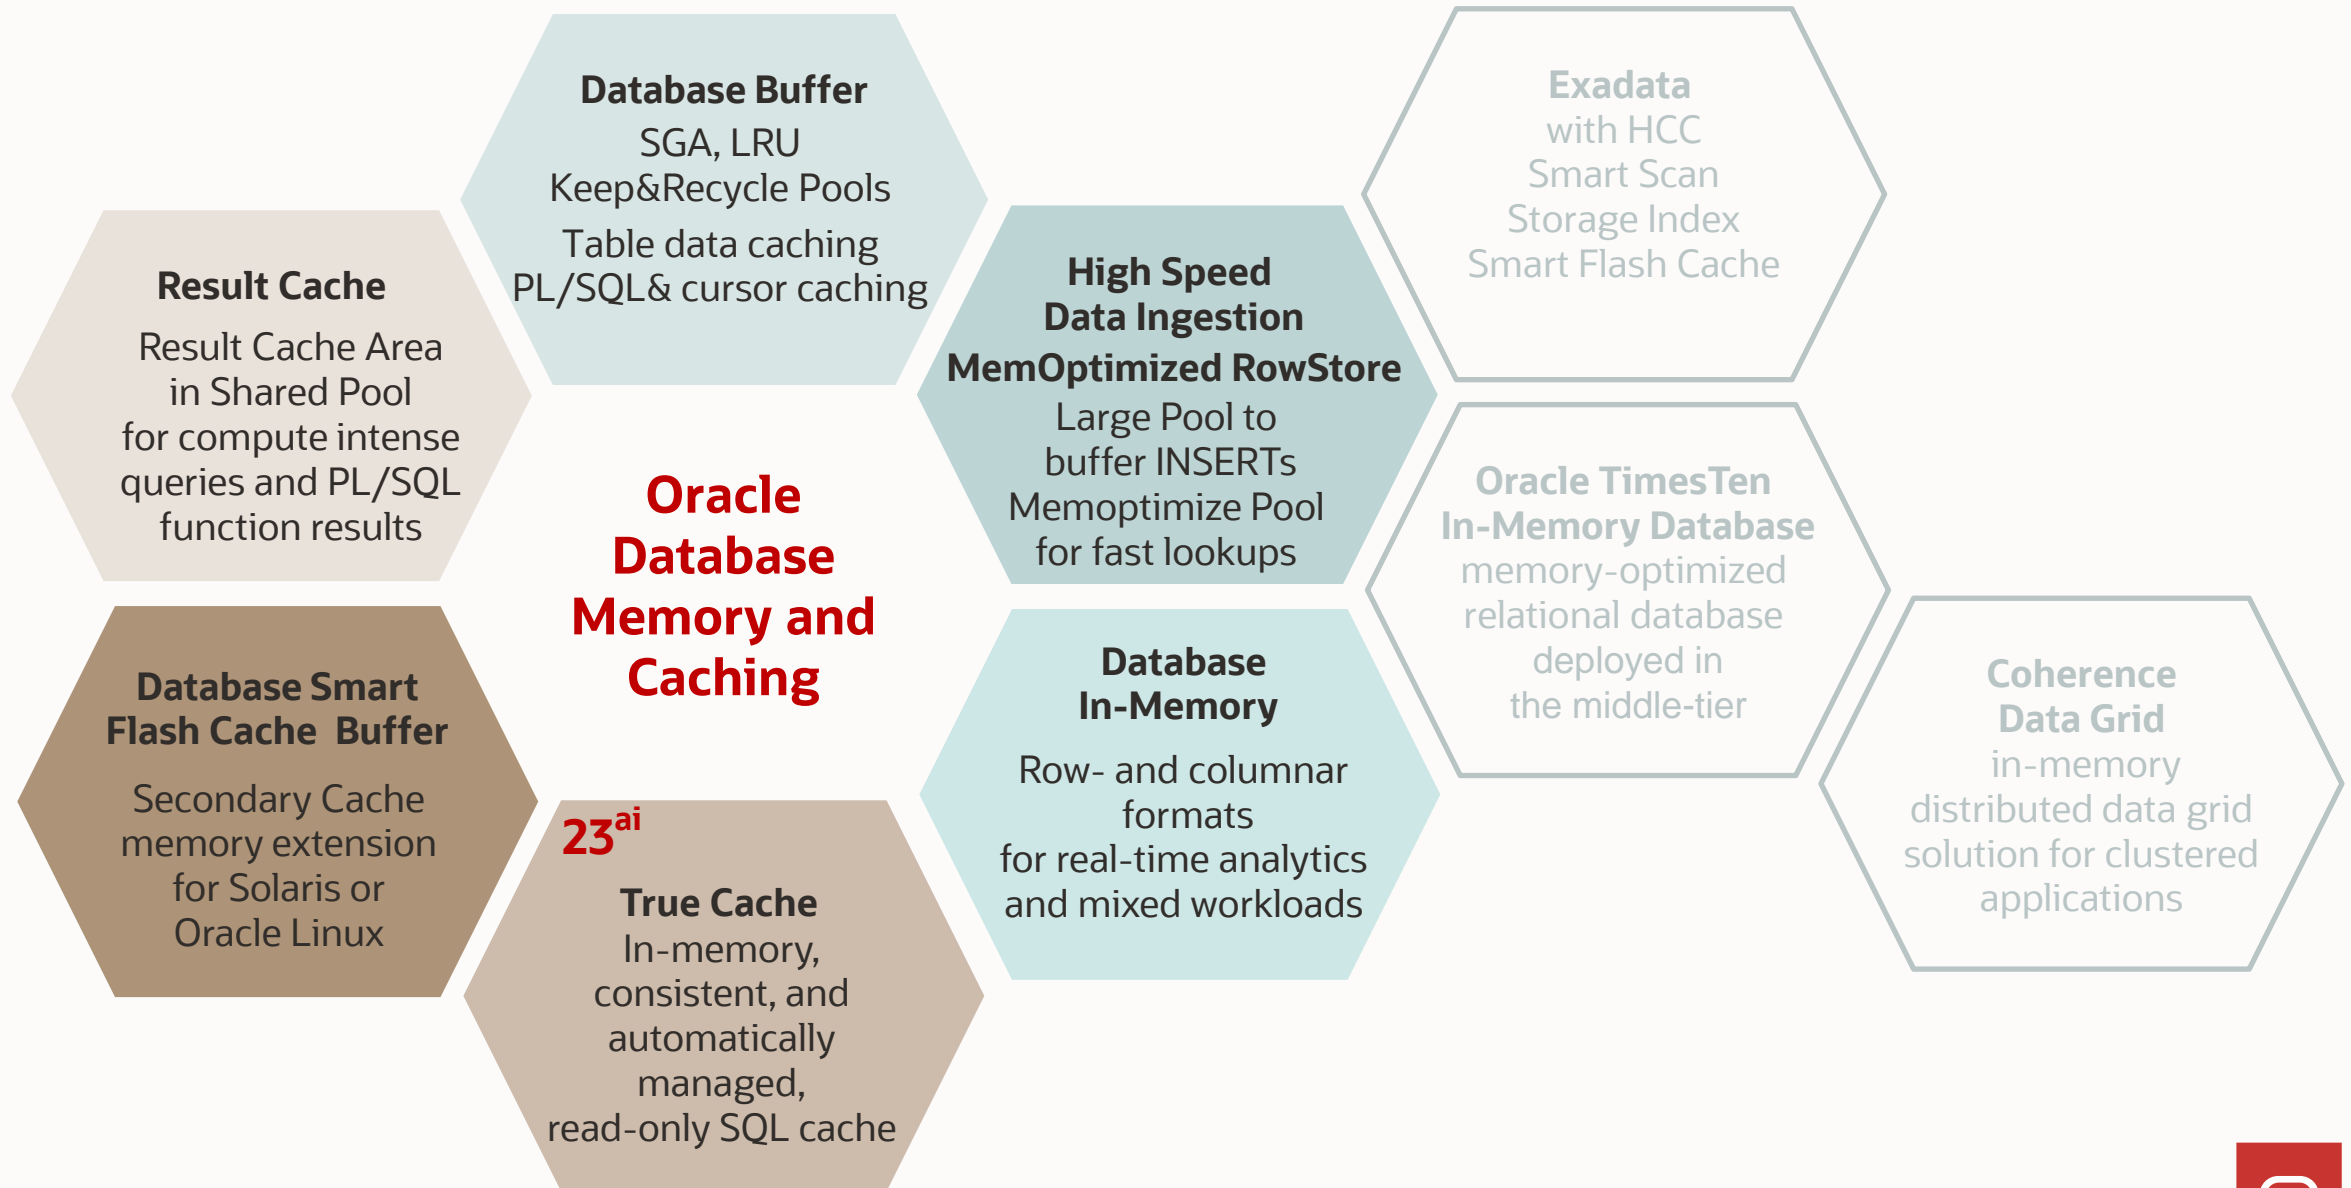
LiveLabs

- [Improve application performance with True Cache](#)

All Links

<https://github.com/oracle-devrel/technology-engineering/tree/main/data-platform/core-converged-db/true-cache>





Videos, Blogs, GitHub from Oracle EMEA Experts and Specialists



YouTube: 23ai Playlist OraTech



GitHub: converged database



Blogs: blogs.oracle.com/coretec

Thank you
